

Assignment 1: Linear regression by using Deep Neural network: Implement Boston housing price prediction problem by Linear regression using Deep Neural network. Use Boston House price prediction dataset

Import necessary libraries

```
import numpy as np # For numerical operations
import pandas as pd # For handling datasets
from sklearn.model_selection import train_test_split # Splitting data into train & test sets
from sklearn.linear_model import LinearRegression # Linear Regression Model
from sklearn.preprocessing import StandardScaler # Standardization of data
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score # Evaluation metrics
```

Importing Keras (for Neural Network)

```
import keras
from keras.models import Sequential # To define a sequential model
from keras.layers import Dense # Fully connected layers
```

Importing Google Colab file handling utility

```
from google.colab import files
```

Uploading and Loading Dataset

```
uploaded = files.upload() # Opens file upload dialogue in Google Colab
boston = pd.read_csv("boston_house_prices.csv") # Reads CSV file into a DataFrame
```

Selecting Features and Target

Selecting 3 input features:

1. LSTAT (Percentage of lower status population)

2. RM (Average number of rooms per dwelling)

3. PTRATIO (Pupil-teacher ratio by town)

```
X = boston[['LSTAT', 'RM', 'PTRATIO']]
```

Target variable: House Price

```
y = boston['PRICE']
```

Splitting the Dataset into Training and Testing Sets

80% of data used for training, 20% for testing

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4)
```

Standardizing the Dataset (Feature Scaling)

Standardization improves model performance by normalizing feature values

```
scaler = StandardScaler() # Initializing StandardScaler
X_train_scaled = scaler.fit_transform(X_train) # Fit and transform training data
X_test_scaled = scaler.transform(X_test) # Transform test data using the same scaler
```

Linear Regression Model

```
lr_model = LinearRegression() # Initializing Linear Regression Model
lr_model.fit(X_train_scaled, y_train) # Training the model using scaled training data
```

Predicting house prices on test data

```
y_pred_lr = lr_model.predict(X_test_scaled)
```

Evaluating Linear Regression Model

```
mse_lr = mean_squared_error(y_test, y_pred_lr) # Mean Squared Error
mae_lr = mean_absolute_error(y_test, y_pred_lr) # Mean Absolute Error
r2_lr = r2_score(y_test, y_pred_lr) # R2 Score (Model accuracy measure)
```

Displaying evaluation metrics

```
print("Linear Regression Model Evaluation:")
print(f"Mean Squared Error: {mse_lr}")
print(f"Mean Absolute Error: {mae_lr}")
print(f"R2 Score: {r2_lr}")
```

Neural Network (ANN) Model

Creating a Deep Learning Model using Keras Sequential API

```
model = Sequential([
    Dense(128, activation='relu', input_dim=3), # Input layer (3 features) & first hidden
    layer (128 neurons)
    Dense(64, activation='relu'), # Second hidden layer with 64 neurons
    Dense(32, activation='relu'), # Third hidden layer with 32 neurons
    Dense(16, activation='relu'), # Fourth hidden layer with 16 neurons
    Dense(1) # Output layer (Predicting a single value - House Price)
])
```

Compiling the model

```
model.compile(optimizer='adam', loss='mse', metrics=['mae'])
```

Optimizer: Adam (Adaptive Learning Rate Optimization)

Loss function: Mean Squared Error (MSE) - Suitable for regression problems

Metric: Mean Absolute Error (MAE) - Helps measure performance

Training the Neural Network

```
history = model.fit(X_train_scaled, y_train, epochs=100, validation_split=0.05,  
verbose=1)
```

Training for 100 epochs

Using 5% of training data as validation set to monitor overfitting

`verbose=1` displays detailed training progress

#Epoch 1/100

#12/12 ————— 4s 26ms/step - loss: 547.8306 - #mae: 21.6359 -

val_loss: 445.7750 - val_mae: 20.1572

#Epoch 2/100

#12/12 ————— 0s 8ms/step - loss: 550.6208 - #mae: 21.6498 -

val_loss: 403.5681 - val_mae: 19.1308

#Epoch 3/100

#12/12 ————— 0s 8ms/step - loss: 433.7596 -

Evaluating the Neural Network Model

```
y_pred_nn = model.predict(X_test_scaled) # Predicting house prices on test data
```

```
mse_nn, mae_nn = model.evaluate(X_test_scaled, y_test) # Evaluating model  
performance
```

Displaying Neural Network Evaluation Metrics

```
print("\nNeural Network Model Evaluation:")
```

```
print(f"Mean Squared Error: {mse_nn}")
```

```
print(f"Mean Absolute Error: {mae_nn}")
```

House Price Prediction for New Data

```
new_data = np.array([[0.1, 10.0, 5.0]])
```

New input values: LSTAT=0.1, RM=10.0, PTRATIO=5.0

```
new_data_scaled = scaler.transform(new_data)
```

Applying the same standardization as training data

Predicting price using trained neural network model

```
prediction = model.predict(new_data_scaled)
```

Displaying the predicted house price

```
print("\nPredicted House Price:", prediction[0][0])
```

#OutPut

1/1 ————— 0s 36ms/step

Predicted House Price: 79.24278